

The Structure of Cross-National Collaboration in OSS

Anonymous Author(s)

Abstract

Open-source software (OSS) development platforms, such as GitHub, expand the potential for cross-national collaboration among developers by lowering the geographic, temporal, and coordination barriers that limited software innovation in the past. However, as shown in previous research, these technological affordances that lower barriers for cross-national collaboration do not uniformly benefit all countries. Using the GitHub Innovation Graph dataset, which aggregates the complete cross-country collaborations among the entire population of GitHub developers, we show the overwhelming dominance of the U.S. in global OSS collaboration and, in stark contrast, the paucity of collaborations among the majority of non-Western country pairs in the data. However, by removing the U.S. and its ties from the collaboration network, we find quantitative evidence of deep-seated religious and cultural affinities, shared colonial histories, and geopolitical factors structuring the collaborations between non-U.S. country pairs. This study highlights the opportunities to develop decentralizing strategies to facilitate new collaborations between developers in non-U.S. countries, thereby fostering the development of novel, innovative solutions. More generally, this study also underscores the importance of contextualizing user behavior and knowledge management in information systems with long-term, macro-social conditions in which these systems are inextricably embedded.¹

Keywords

open-source software, international collaboration, cultural homophily, network analysis, GitHub, exponential random graph models, ERGM, structural equivalence, world systems theory, civilizations, geopolitics, hierarchical clustering, core-periphery, software collaboration, cross-border collaboration, block models

ACM Reference Format:

Anonymous Author(s). 2025. The Structure of Cross-National Collaboration in OSS. In *Proceedings of The 34th ACM Conference on Information and Knowledge Management (CIKM'25)*. ACM, New York, NY, USA, 10 pages. <https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

1 Introduction

How do countries collaborate in open-source software development? Throughout history, dense economic trade relations in commodities, goods, services, and labor tended to promote international

peace whereas sparse connections and fragmentation heightened the probability of war [13]. These international connections among countries have been subject to a myriad of social, economic, and political forces with shifting influences over time. For example, Western capitalist versus Eastern communist countries during the Cold War era formed connections along ideological fault lines while the post-Cold War divisions have been structured in part by deep-seated culture and religion [12]. Their significance has been shown to persist in online social interactions among millions of email and social media users, despite the friction-less connectivity that these communication technologies enabled [22].

While studies show the persistence of such ideological, cultural, and geopolitical forces on cross-national online communication, it is an open question whether cross-national collaborations with a clearer instrumental focus, such as those observed in open-source software development, also form under the influence of such macro-social forces. Presumably, these bottom-up collaborations between millions of software developers across countries should appear oblivious to these social forces insofar as the online collaboration platform helps overcome spatio-temporal barriers of collaboration in the creation of technological and economic value. Indeed, previous studies of open-source software development have shown that cross-country collaborations largely exhibit such instrumental rationality, reflecting the economic dominance of the U.S. and, to a lesser extent, of a handful of technologically advanced European countries. Although this heavy reliance on the U.S. is taken for granted, it is puzzling from the vantage point of the potential for frictionless and untethered collaborations that OSS platforms afford.

Based on a recent cross-country collaboration dataset released by GitHub, we construct a network of OSS collaborations between countries and study the patterns hidden under the global dominance of the United States. We reveal the structural markers of international OSS collaboration that indicate recent geopolitical shocks as well as deep-seated, macro-cultural fault lines that have been observed in cross-border email and Twitter communications [22]. Despite the common belief that OSS collaboration can overcome geographical barriers [4], we observe cultural affinities (i.e., cultural homophily) and historical influences (shared colonial past) that create social boundaries in inter-country collaborations.

Our paper makes two primary contributions. First, we contribute to OSS research by proposing to remove the dominant U.S. for exploring otherwise hidden structures that provide insight. Second, we contribute to data mining and knowledge management in collaborative information systems by demonstrating the deep insights that can be gained by considering social processes that unfold at historical time scales, an approach that has been neglected in the literature.

¹The code with anonymization is available at: <https://zenodo.org/doi/10.5281/zenodo.10521115>

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
CIKM'25, Seoul, Korea

© 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM ISBN 978-x-xxxx-xxxx-x/YYYY/MM
<https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

2 Related Work

2.1 Structure of Global Interdependence

A prevailing view on international relations is that the interdependence among nations lead to a more integrated world. Dense connections between countries through economic exchange of goods, services, and labor create strong economic interdependence that inhibit war and lead to prolonged prosperity [13]. From this perspective, the OSS software development collaboration ties that emerge between countries through the uncoordinated activities of the hundreds of thousands of individual software developers are a potentially important aspect to consider for understanding international economic interdependence in today's information economy. Furthermore, the structure of OSS collaboration is an important consideration in terms of technological innovation, since these mutual collaborations across borders directly affect the extent of joint-innovations, and their global diffusion potential [1]. As such, gaining a deeper understanding of the principles that govern bottom-up open-source software development across national borders can reveal novel collaboration opportunities for innovation.

2.2 Hierarchy in Global Open-Source Collaboration

A widespread belief about OSS development is that it can remove geographic barriers to enable global collaboration & innovation at an unprecedented rate. However, the reality is that OSS development also clusters geographically [24]. Earlier studies in Linux [5], SourceForge [7, 25], and GitHub [23] consistently indicate that OSS development is dominated by developers from the U.S., the EU, as illustrated in striking visualizations [10]. Even within this concentration of OSS activity and collaborations among Western countries, studies also highlight the extreme hierarchy and asymmetry between the U.S. and other European nations, where software projects led by U.S. developers attracts exceptional levels of international contributions relative to the contributions that U.S. developers tend to make in non-U.S. projects Takhteyev and Hiltz [23]. Although, some recent studies [19, 26] show that geographic concentration in development activity has been gradually decreasing since the 2000s (i.e., participation of non-U.S. and non-EU countries), the projects based in the dominant countries tend to reject code contributions from less developed countries [6, 16, 17].

2.3 Deep-Seated Cultural and Historical Forces

With the unmistakable U.S. dominance in open-source software collaboration, it is difficult not to lose sight of the macro-historical forces underlying international politics, trade, and conflict. For example, the ideological schisms of the Cold War era had a profound, enduring impact on global interactions.

Samuel Huntington Huntington [12] posited that the fault lines in post-Cold War international relations will shift from ideological to deep-seated cultural and religious divides. These divides proved to be important in the formation of distinct blocs and networks, influencing not only political alliances but also scientific, technological, and industrial collaborations. This historical context sets the stage for understanding the paradoxical role of modern communication technologies in global connectivity, as explored in [22].

Challenging the notion that advancements in communication technologies inherently lead to a more connected world, [22] reveals that despite the global reach of digital communication platforms, interactions often reflect and sometimes exacerbate existing cultural and national divides. The findings suggest a preference for in-group communication, even in digital spaces, indicating that technological tools alone are insufficient for narrowing cultural gaps.

These works collectively paint a picture of global interdependence that is far more complex than the OSS narrative of digital integration and untethered connectivity might suggest. They illustrate how economic, ideological, and technological factors, while having the potential to bring nations closer, also maintain or even widen existing divides. Such a nuanced understanding of global interdependence is crucial for considering the dynamics of international collaboration, particularly in fields like open-source software development. Hence, moving beyond general-purpose communications, we test the effects of broad cultural/religious commonalities on the structure of the specialized, goal-driven collaboration network in GitHub.

3 Methods

3.1 Data

In this study, we use the publicly available GitHub Innovation Graph dataset² which provides the complete quarterly country-to-country *Git push*³ volume data from Q1, 2020 to Q2, 2023. The dataset infers developer locations based on the IP addresses associated with a developer's Git pushes.⁴

We study the aggregate collaborations across all quarters by creating a weighted, directed graph consisting of countries as nodes and their aggregate Git pushes from one country to another as weighted edges between them. However, this raw graph is heavily skewed with weights between major countries that are an order of magnitude larger than the weights between smaller, developing countries. Given our goal of delineating the multi-scale international collaborations while filtering out the possible noise observed in the low-weight edges, we first normalize the weight of each edge by the total outgoing weight of its source, then retain only the edges that are more likely to have structural significance, using the disparity filter for uncovering the network backbone [20].

The disparity filter examines the weights of a node's outgoing ties and retains the edges with relatively large weights that are statistically unlikely if the edge weights of that node had been distributed randomly. The filter is customizable with an alpha value at which statistical significance is defined. Choosing a low alpha value results in more aggressive edge filtering. In contrast to uniform thresholding (i.e., filtering out any edge with weight below a fixed threshold), the disparity filter adapts to the local structure of the network, preserving the heterogeneity of connections.

²<https://innovationgraph.github.com/>

³The GitHub platform uses the Git version control system. In Git parlance, contributions to a source code repository are packaged as "commits" that get "pushed" to that repository.

⁴<https://github.com/github/innovationgraph/blob/main/docs/datasheet.md>

3.2 Hierarchical Clustering

To characterize the hierarchical structure of the global OSS collaboration network and identify structural equivalence classes (block models), we employ hierarchical clustering analysis [27]. This approach groups countries based on the similarity of their collaboration patterns rather than their direct connections, revealing the underlying positional structure of the network consistent with world systems theory [21].

We construct structural equivalence by measuring the Euclidean distance between countries' collaboration profiles. For each pair of countries i and j , we compute:

$$d_{ij} = \|\mathbf{a}_i - \mathbf{a}_j\|_2 = \sqrt{\sum_{k=1}^n (A_{ik} - A_{jk})^2} \quad (1)$$

where $\mathbf{a}_i$ denotes the i th row of the adjacency matrix A , representing country i 's collaboration profile, and n is the number of countries. Countries with similar distance values have comparable patterns of collaboration with other countries, indicating structural equivalence regardless of whether they directly collaborate with each other [27].

Using the complete linkage method, we perform hierarchical agglomerative clustering on the distance matrix. The optimal number of clusters is determined by maximizing the Calinski-Harabasz score, which evaluates cluster separation and compactness [3]. Formally, the Calinski-Harabasz (CH) index for k clusters and N countries is defined as:

$$\text{CH}(k) = \frac{\text{BCSS}/(k-1)}{\text{WCSS}/(N-k)} \quad (2)$$

where the between-cluster sum of squares (BCSS) and within-cluster sum of squares (WCSS) are computed as:

$$\text{BCSS} = \sum_{i=1}^k n_i \|\mathbf{c}_i - \bar{\mathbf{a}}\|_2^2, \quad \text{WCSS} = \sum_{i=1}^k \sum_{\mathbf{a} \in C_i} \|\mathbf{a} - \mathbf{c}_i\|_2^2 \quad (3)$$

Here, n_i is the number of countries in cluster i , $\mathbf{c}_i$ is the centroid of cluster i , and $\bar{\mathbf{a}}$ is the global centroid of all collaboration profiles. All distances are measured in the same Euclidean space as defined above. A higher CH score indicates more distinct and compact clustering. We evaluated clustering solutions across threshold values from 20 to 150 with unit intervals, selecting the configuration with the highest CH score.

The resulting block model partition identifies structural positions within the global collaboration hierarchy [21]. Countries in the same block share similar roles in the network: they may not collaborate directly with each other, but they maintain comparable patterns of relationships with countries in other blocks. This structural equivalence approach reveals the core-periphery dynamics where the core countries (typically advanced economies) receive disproportionate collaboration from semi-peripheral and peripheral countries, while peripheral countries exhibit sparse internal connectivity despite occupying similar structural positions [21].

This block modeling approach differs fundamentally from community detection methods that identify densely connected groups. Instead, it reveals the positional structure of countries based on their equivalent roles in channeling international OSS collaboration

flows, providing insight into the hierarchical organization of global technological networks consistent with the theoretical predictions of world systems [21].

3.3 Exponential Random Graph Models

To quantitatively assess the influence of cultural factors on cross-national OSS collaboration patterns while accounting for the complex interdependencies inherent in network data, we employ Exponential Random Graph Models (ERGMs) [18]. Traditional statistical approaches that assume independence among observations are inappropriate for network analysis, as the formation of one collaborative tie may directly influence the probability of forming additional ties within the same network structure [27]. ERGMs address this fundamental limitation by explicitly modeling the conditional dependence structure of network formation processes.

ERGMs belong to the exponential family of probability distributions and model the likelihood of observing a particular network configuration $\mathbf{y}$ as:

$$P_{\theta}(\mathbf{Y} = \mathbf{y}) = \frac{1}{\kappa(\theta)} \exp \{ \theta^T \mathbf{g}(\mathbf{y}) \} \quad (4)$$

where $\theta = (\theta_1, \theta_2, \dots, \theta_p)^T$ represents the vector of model parameters, $\mathbf{g}(\mathbf{y}) = (g_1(\mathbf{y}), \dots, g_p(\mathbf{y}))^T$ denotes the vector of sufficient statistics capturing various network configurations, and $\kappa(\theta)$ is the normalizing constant:

$$\kappa(\theta) = \sum_{\mathbf{y}^* \in \mathcal{Y}} \exp \{ \theta^T \mathbf{g}(\mathbf{y}^*) \} \quad (5)$$

where $\mathcal{Y}$ represents the sample space of all possible networks on n nodes.

The conditional log-odds of edge formation between nodes i and j , given the rest of the network $\mathbf{Y}_{ij}^c$, can be expressed as:

$$\text{logit}(P(Y_{ij} = 1 | \mathbf{Y}_{ij}^c)) = \theta^T \delta(\mathbf{y})_{ij} \quad (6)$$

where the sum is over all configurations A that contain Y_{ij} , and $\delta_A(\mathbf{y})$ is the change statistic representing the change in the value of the network statistic $g_A(\mathbf{y})$ when y_{ij} changes from 0 to 1.

Network Statistics and Model Terms. Our ERGM specification incorporates several theoretically motivated network statistics to test hypotheses regarding cultural homophily, reciprocity, and structural patterns in international OSS collaboration [11]:

Edges Term: The fundamental density statistic controls for the baseline propensity of collaboration:

$$g_{\text{edges}}(\mathbf{y}) = \sum_{i,j} y_{ij} \quad (7)$$

The parameter θ_{edges} represents the log-odds of collaboration in the absence of other effects.

Nodematch Terms: To test Huntington's cultural homophily hypothesis [12], we employ nodematch statistics that quantify the tendency for countries sharing civilization membership to exhibit higher collaboration rates [14]. For civilization attribute C , the differential nodematch statistic for civilization k is:

$$g_{\text{nodematch},k}(\mathbf{y}) = \sum_{i,j} y_{ij} \mathbf{1}(C_i = C_j = k) \quad (8)$$

where $\mathbf{1}(C_i = C_j = k)$ is an indicator function equal to 1 when both countries i and j belong to civilization k . This specification allows for civilization-specific homophily coefficients, enabling us to test whether different cultural groups exhibit varying degrees of preference for internal collaboration.

For structural equivalence effects, we also test block model homophily using:

$$g_{\text{block model}}(\mathbf{y}) = \sum_{i,j} y_{ij} \mathbf{1}(B_i = B_j) \quad (9)$$

where B_i represents the block model position of country i derived from hierarchical clustering of structural equivalence patterns.

Mutual Terms: For directed networks, reciprocity effects are captured through mutual dyad statistics [27]:

$$g_{\text{mutual}}(\mathbf{y}) = \sum_{i,j} y_{ij} y_{ji} \quad (10)$$

A positive coefficient θ_{mutual} indicates that countries receiving OSS contributions demonstrate a higher propensity to reciprocate, suggesting sustained collaborative relationships.

Baseline Model Specification and Estimation. Our baseline model specification evaluates the structure of the collaboration network while controlling for density and reciprocity effects:

$$\Pr(\mathbf{Y} = \mathbf{y}) = \frac{1}{k(\boldsymbol{\theta})} \exp \{ \theta_{\text{edges}} L(\mathbf{y}) + \theta_{\text{mutual}} M(\mathbf{y}) \} \quad (11)$$

where $L(\mathbf{y}) = \sum_{i,j} y_{ij}$ is the number of edges, $M(\mathbf{y}) = \sum_{i,j} y_{ij} y_{ji}$ is the number of mutual dyads.

Parameter estimation is conducted using Monte Carlo Maximum Likelihood Estimation (MCMC-MLE) as implemented in the `ergm` package [9]. The MCMC algorithm approximates the intractable normalizing constant through iterative simulation from the model distribution. We employ a burn-in period of 40,000 iterations with sampling intervals of 10,000 iterations to ensure adequate mixing and convergence. Model adequacy is assessed through goodness-of-fit diagnostics that compare observed network statistics to distributions generated from the fitted model.

3.4 Node2Vec

We test the robustness of both the hierarchical clustering of countries in structural equivalent positions and the direct collaboration ties based on cultural similarity in a unified node embedding framework, node2vec. Node2vec learns continuous vector representations of nodes by performing biased random walks controlled by two hyperparameters [8]. The return parameter p controls the likelihood of revisiting nodes, while the in-out parameter q controls whether walks explore locally or globally.

Different parameter combinations capture distinct network properties:

Homophily ($p = 4, q = 0.25$): High p and low q encourage local exploration, capturing direct collaboration patterns between similar countries.

Structural equivalence ($p = 0.25, q = 4$): Low p and high q emphasize global network roles, capturing countries that play similar structural positions (e.g., hub countries) regardless of direct connections.

We selected these configurations to test whether colonial patterns manifest through direct collaboration preferences (homophily) or similar structural roles in the global network (structural equivalence). This allows systematic examination of different mechanisms through which historical relationships may persist in modern OSS collaboration networks.

3.5 Colonial Dimension Analysis

Building on the node2vec embeddings described above, we develop a colonial dimension methodology to investigate whether historical colonial relationships persist in the OSS collaboration network structure. This approach tests whether countries exhibit systematic proximity to their historical colonizers when positioned along a theoretically motivated vector space derived from colonial relationships.

Using the node2vec embeddings, we construct a "colonized dimension" by computing the average vector from three major historical colonial relationships. Using colonial pairs India-Britain (IN-GB), Vietnam-France (VN-FR), and Mexico-Spain (MX-ES), the colonized dimension is defined as:

$$CD = \frac{(v_{IN} - v_{GB}) + (v_{VN} - v_{FR}) + (v_{MX} - v_{ES})}{3} \quad (12)$$

where v_i represents the node2vec embedding for country i computed using the parameters specified in Section 3.4. This dimension captures the average directional vector from former colonizers toward their colonies, theoretically representing a "decolonized" direction in the embedding space.

We test this dimension by projecting countries of interest along the colonized vector ($v_i + CD$) and measuring cosine similarity to potential colonizers. Test countries and regions include Ukraine, Poland, South Korea, Taiwan, Argentina, Brazil, Canada, Senegal, Turkey, United States, and China. The analysis spans the same four configurations described in Section 3.4 along with two more balanced parameter settings. The final settings are three node2vec parameter settings (structural equivalence $p=0.25, q=4$; homophily $p=4, q=0.25$; balanced $p=1, q=1$) under two network conditions (with and without U.S.).

4 Results

4.1 Hierarchical Structure of Global OSS Collaboration

The hierarchical clustering analysis reveals four distinct structural positions in the global OSS collaboration network. The dendrogram in Figure 1 demonstrates how countries naturally cluster based on their structural equivalence patterns, with the resulting blocks reflecting real-world geopolitical hierarchies.

The dendrogram shows distinct clustering patterns across all country groups. Among the periphery countries (shown in orange), we observe clear regional groupings where African countries cluster together, Latin American countries form another coherent cluster, and small island nations and developing economies occupy separate branches. The dendrogram structure indicates that despite their geographic dispersion, these countries share similar structural positions in the global OSS network—they predominantly contribute




Figure 1: Hierarchical clustering dendrogram of all countries in the global OSS collaboration network. The blue lines highlight the separation of the U.S. from other countries, reflecting its unique structural position. Countries are colored by their structural equivalence blocks: periphery (orange), core (red), semi-periphery (green), and the U.S. as a separate block. The optimal cutting points for the four-block partition were determined by maximizing the Calinski-Harabasz score, revealing the hierarchical organization of the global OSS network.

to projects in core countries while receiving minimal collaboration from other periphery nations.

The dendrogram also illustrates the hierarchical relationships among core and semi-periphery countries. The core countries (shown in red) include major Western European nations (Netherlands, Germany, France, United Kingdom), along with economically advanced countries like Canada and Spain. These countries occupy similar structural positions as regional hubs that attract substantial collaboration from their respective spheres of influence. The semi-periphery countries (shown in green) display more heterogeneous clustering patterns, reflecting their intermediate position between core and periphery. Notable clusters include Eastern European countries (Poland, Czechia, Romania), Asian emerging economies (Taiwan, South Korea, Singapore), and other middle-income countries that serve as bridges between the global core and their regional peripheries.

The blue lines in the dendrogram highlight the unique structural position of the United States, which separates early in the clustering process, confirming its exceptional role in the global OSS collaboration network that we analyze in detail below.

The global collaboration pattern exhibits a strict hierarchy among four groups of countries that assume structurally equivalent positions, or “blocks”, obtained from a block model analysis. Figure 2 plots the asymmetry in collaborations among these four blocks – core, semi-periphery, periphery blocks and the U.S. as its own block. The cells represent a column block’s relative collaboration volume directed to a row block (e.g., the collaboration volume from developers in peripheral countries on U.S. projects is 233.78% of the U.S. developers’ collaboration volume on the projects of the peripheral countries). The volume of collaborations directed between clusters are highly asymmetric and transitive – as shown in Figure 2, the core countries consistently receive disproportionately more collaborations from the semi-periphery and the periphery countries while the semi-periphery countries receive out-sized collaborations from the periphery countries. This perfectly transitive structure remains highly stable across quarters with the U.S. at the apex, receiving disproportionate collaborations from the other three blocks.

Indeed, the strong centripetal force of the U.S. is visually apparent in Figure 3a, which displays a subset of countries that appear in Samuel Huntington’s civilization classification, based on a force directed layout (Force Atlas 2 in Gephi [2]). Node colors represent countries within densely knit communities of strong collaboration ties, as labeled by the Louvain community detection algorithm. Consistent with previous studies that report the exceptional U.S. prominence in OSS development, it looms large at the center, so much so that over 60% of the countries are labeled into the same community as the U.S., gravitating en masse, thereby blurring the more subtle boundaries that we present below.

However, once the U.S. is removed from the graph, the overall hierarchy obtained from the aforementioned block model is more apparent. In Figure 4b with the U.S. removed, the node colors are based on the structurally equivalent blocks of the core (red), semi-periphery (green), and periphery (purple) countries. In the absence of the U.S., the large core countries (i.e., Great Britain, Germany, France, Russia, and India) are spatially dispersed, with their respective semi-periphery countries closely positioned.




Figure 2: Relative collaboration tie strength between core, semi-periphery, and periphery country positions


4.2 Cultural Homophily and Clustering

We present four ERGMs that incrementally evaluate the influence of structural equivalence and cultural homophily on cross-national OSS collaboration. Each model builds on the baseline network structure while testing additional hypotheses, emphasizing cultural affinity and positional equivalence as sources of clustering.


Model 1: Baseline Network Structure. The baseline model includes only the number of edges and mutual collaboration ties. The results in Table 1 show a strong and significant mutuality effect (2.294, $p < 0.001$), indicating that dyads with bilateral contributions are substantially more likely to form. However, it is a stretch to interpret mutual as reciprocity in the sense of quid-pro-quo between countries. Since ties are aggregates of individual behaviors, contributors are unlikely to coordinate country-to-country exchanges. Rather, the mutual effect likely reflects broader similarity patterns between countries that are not captured by other covariates.

Model 2: Structural Equivalence (Same Block). Model 2 adds a nodematch term for countries in the same structural equivalence block derived via hierarchical clustering. Table 1 shows that the coefficient is positive but only marginally significant (0.165, $p < 0.1$). According to world system theory [21], countries occupying peripheral positions should exhibit low internal connectivity, with stronger outward links to the core. Hence, the limited magnitude and marginal significance of the same-block effect is not a discouraging result, but consistent with the expected hierarchical structure of global collaboration.


Model 3: Cultural Homophily (Civilizations). Model 3 evaluates the influence of shared cultural affiliation by including nodematch terms for Huntington’s civilization categories. Table 2 shows that most civilization terms are strongly significant and positive. The **Sinic civilization** exhibits the strongest internal collaboration



(a) Nodes including the U.S., colored by modularity class



(b) Nodes excluding the U.S., colored by block model positions (red: core, green: semi-periphery, purple: periphery)



(c) K-means clustering (7 clusters) on node2vec embeddings emphasizing structural equivalence (p=0.25, q=4)

Figure 3: Force Atlas layout of the GitHub collaboration graph including the U.S. and colored by modularity class (a), without the US colored by block model positions (b), and k-means clustering applied to node2vec embeddings that emphasize structural equivalence patterns (c). Node size and edge thickness are proportional to weighted nodal degree and edge weight, respectively.

Table 1: ERGM Models 1 & 2: Baseline and Structural Equivalence

Term	Model 1	Model 2
Edges	−3.206***	−3.263***
Mutual	2.294***	2.283***
Nodematch.SameBlock	−	0.165*

*** $p < 0.001$, * $p < 0.1$ **Table 2: ERGM Models 3 & 4: Cultural Homophily**

Term	Model 3	Model 4
Edges	−3.518***	−3.490***
Mutual	1.718***	1.730***
Nodematch.SameBlock	−	−0.114
Nodematch.African	0.775**	0.841**
Nodematch.Buddhist	1.902***	1.939***
Nodematch.Islamic	0.313	0.320
Nodematch.LatinAmerican	1.657***	1.690***
Nodematch.Orthodox	1.725***	1.707***
Nodematch.Sinic	2.344***	2.384***
Nodematch.Western	1.604***	1.630***

*** $p < 0.001$, ** $p < 0.01$

tendency (2.344, $p < 0.001$), followed by **Orthodox** (1.725, $p < 0.001$) and **Latin American** (1.657, $p < 0.001$) civilizations. In contrast, **Islamic civilization** shows no significant homophily effect (0.313, n.s.), while **African civilization** demonstrates weaker but significant clustering (0.775, $p < 0.01$). These results are visually apparent in Figure 4a where the African countries and Islamic countries are rather scattered compared to other tightly clustered non-Western civilizations (e.g., Latin American, Orthodox).

Model 4: Combined Civilizations and the Same Block. Model 4 incorporates both civilization categories and the positional equivalence of the same block. As shown in Table 2, the civilization coefficients remain robust and significant, while the same block term becomes negative and nonsignificant (−0.114, n.s.). This reinforces the interpretation that structural equivalence captures countries with similar positions in the global system, not necessarily countries with stronger bilateral collaborations. Rather than predicting direct ties, block membership is useful for identifying countries that share similar patterns of international collaboration.

Two countries in the same block may or may not have direct ties, because being in the same block simply means that they have similar connection patterns to others (i.e., structural equivalence), not that they should have a higher propensity to form direct ties with each other. According to world systems theory [21], the periphery block should have low internal connectivity, with high interconnections within the core block, fewer within the semiperiphery, and sparse connections within the periphery itself. This is a great case for using block models: their strength lies not in predicting direct connections, but in identifying shared relational positions in a stratified global hierarchy. The contrast with Model 2 (Same Block only: 0.165, $p < 0.1$) further demonstrates that structural equivalence alone does not provide strong explanatory power. In




**(a) Colored by Huntington's civilization categories****(b) Colored by modularity class****(c) K-means clustering (7 clusters) on node2vec embeddings emphasizing homophily ($p=4$, $q=0.25$)**

Figure 4: Force Atlas layout of the GitHub collaboration graph without the U.S., colored by Samuel Huntington's civilization labels (a), modularity class (b), and k-means clustering applied to node2vec embeddings that emphasize homophily patterns (c). Node size and edge thickness are proportional to weighted nodal degree and edge weight, respectively.

contrast, Model 4 shows that cultural homophily, operationalized through shared civilizational identity, offers more consistent and interpretable insight into collaboration tendencies in global OSS production.

4.3 Imperialism and Geopolitics

The modularity coloring in Figure 4b hints at historical factors that continue to structure OSS collaborations. Of particular note, the

modules rather accurately identify common colonial history with, for example, the former British colonies located in the Northwest, neighbored by the former French colonies in the North, and by the former Spanish colonies to the Southwest. The former Soviet bloc countries are also grouped as one community.

To quantitatively test the persistence of these colonial relationships, we employ the colonial dimension methodology described in Section 3.5. This analysis provides direct evidence for historical patterns embedded within the collaboration network structure. **Japan consistently emerges as the closest potential colonizer for South Korea and Taiwan** (4/6 configurations), accurately reflecting historical Japanese colonization. **Spain matches Argentina in 3/6 configurations** and **Portugal matches Brazil in 2/6 configurations**, confirming Iberian colonial legacies persist in Latin American OSS patterns. Beyond formal colonialism, **Russia frequently appears for Ukraine and Poland**, reflecting Soviet sphere influence rather than traditional colonial relationships.

However, contemporary superpowers demonstrate exceptional resilience to historical projection. Both the **United States and China consistently exhibit highest similarity to themselves** (>0.99 cosine similarity), indicating these nations maintain such dominant structural positions that even colonial dimension adjustments cannot overcome their network centrality.

The past imperial connections between the colonizers and the colonized are also detected through node embeddings. We generated the classic node2vec embeddings with hyperparameters favoring structural equivalence to highlight the structural role equivalences of the colonizers (France, United Kingdom, Spain) as well as super powers in the cold war era (Russia) on the one hand and their subordinate countries on the other. Specifically, the embeddings were generated for different combinations of return parameter p and in-out parameter q , with values $[0.25, 0.5, 1, 2, 4]$ for both p and q . For each combination of p and q , we generated node2vec embeddings, standardized them using StandardScaler, and ran k-means clustering with 7 clusters using Euclidean distance as the similarity metric. Figure 4c shows the community structure based on homophily parameters ($p=4, q=0.25$) and figure 3c shows the community structure based on structural equivalence parameters ($p=0.25, q=4$) with the node colors indicating community membership.

The comparison across different analytical approaches reveals both the robustness and underlying complexity of these historical patterns. Both hierarchical block modeling and k-means clustering based on node2vec with structural equivalence parameters consistently identify core geopolitical clusters—Western European powers (GB, DE, FR, NL), Eastern European countries (RU, PL, CZ), and major Asian economies (JP, KR, IN)—demonstrating that fundamental structural roles persist across methodologies. However, the node2vec approach fragments the interpretable three-tier hierarchy into more granular communities that capture subtle structural similarities, such as separating China (CN) and India (IN) into distinct groups, reflecting nuanced collaboration patterns that lack clear theoretical interpretation. Similarly, while both modularity-based community detection and k-means clustering based on node2vec with homophily parameters identify core colonial spheres—the British cluster (GB, CA, AU, NZ), French Francophone networks, and Spanish-Latin American connections—they diverge in their

treatment of peripheral countries. The modularity approach maintains clear colonial alignments where African countries cluster with their respective former colonial powers, while the node2vec approach consolidates many African countries into separate pan-African communities, suggesting the algorithm detects contemporary South-South collaboration patterns that transcend colonial relationships. These methodological differences illustrate the fundamental trade-off between algorithmic granularity and interpretable hierarchy, demonstrating how homophily and structural equivalence operate at multiple scales within both historical networks and emerging regional collaborations.

Taken together, the apparent organization of collaboration around common colonial history (and cold war hierarchies in the Soviet block) may partly explain why countries from the Islamic and African cultures are less visually clustered within their own culture, but more concentrated around the former imperialist countries. The economic dependencies that date back to the colonial era continue to influence today's knowledge economy that is supported in part by OSS development.

5 Discussion

Extending previous studies that identify the geographic locus of concentrated development in OSS, this study dissects the structure of the cross-national collaboration network on Github. As reported in the literature, we observe a strictly hierarchical collaboration structure with the U.S. as the strongest gravitational force at the center of the global collaboration network, followed by a core group of technologically advanced countries heavily contributing to U.S.-based projects. This core group, in turn, hosts projects to which the semi-periphery and periphery countries heavily contribute. This clear hierarchy in terms of core, semi-periphery, and periphery countries is reminiscent of the trade network structure that social network analysts have discovered from input-output tables in the 1970s [21].

Furthermore, these international collaborations also appear to reflect long-term cultural affinities and shared historical experiences between countries. By triangulating the results from network visualizations, block models, exponential random graph models, and node embeddings, we demonstrated the robust influence of deep-seated cultural and religious similarities (i.e., civilizations) as well as long-lasting geopolitical power structures from the past (i.e., colonial power-dependencies and cold-war super powers).

5.1 Implications and Future Directions

Our findings offer several key insights with broad implications:

5.1.1 Influence on Commercial Software Development Practices. Our results indicate that geopolitical factors have a significant impact on OSS contributions. This calls for a reevaluation of global outsourcing and collaboration strategies in commercial software development. Future research should focus on developing more agile and adaptable models for global software teamwork.

5.1.2 Overcoming Barriers in OSS Collaboration. Identifying and addressing both inter-community and intra-community barriers in OSS collaboration presents a significant challenge. Future initiatives might focus on developing enhanced recommender systems for

OSS platforms, incorporating cultural tags to foster more effective cross-cultural collaborations that are conducive to innovation [15]. Additionally, there is an opportunity to develop features within platforms like GitHub to streamline collaboration within culturally similar groups, potentially transforming these platforms into more effective multi-hub network structures.

5.1.3 Broader Research Horizons. Future research might also explore the socio-political dimensions of OSS collaboration, examining how global political shifts and policy changes influence the dynamics of these networks. Understanding these broader impacts could provide valuable insights into the evolving nature of global software development and collaboration.

6 Conclusion

The research also documents the possible impact of geopolitical events and platform censorship policies on OSS activities, such as China's decreasing contributions possibly due to censorship and the effects of the Russia-Ukraine conflict. Additionally, it brings to light the enduring influence of colonial history on current OSS collaborations.

Implications of this study are far-reaching, suggesting the need for resilient OSS community frameworks, adaptable commercial software development strategies, and tools to overcome both inter and intra-community collaboration barriers. Future research directions include exploring the socio-political dimensions of OSS collaboration and how global changes impact these networks. While the findings of this study are informative, they are to be evaluated with the following limitations with respect to data. Specifically, the reliance on IP-based activity statistics introduces potential biases, particularly in countries with prevalent VPN usage, like China. This could lead to distorted representations of actual OSS activity. While we anticipate these distortions are limited in scope, their impact on the study's findings should not be underestimated. Future research might explore alternative data collection methods to mitigate this limitation.

Generative AI Disclosure

No generative AI tools were used in any stage of this research, including data collection, analysis, code development, or manuscript preparation.

References

- [1] Tufool Alnuaimi, Jasjit Singh, and Gerard George. 2012. Not with my own: long-term effects of cross-country collaboration on subsidiary innovation in emerging economies versus advanced economies. *Journal of Economic Geography* 12, 5 (2012), 943–968. <http://www.jstor.org/stable/26158628>
- [2] Mathieu Bastian, Sebastien Heymann, and Mathieu Jacomy. 2009. Gephi: An Open Source Software for Exploring and Manipulating Networks. <http://www.aaai.org/ocs/index.php/ICWSM/09/paper/view/154>
- [3] T. Caliński and J. Harabasz and. 1974. A dendrite method for cluster analysis. *Communications in Statistics* 3, 1 (1974), 1–27. doi:10.1080/03610927408827101 arXiv:https://www.tandfonline.com/doi/pdf/10.1080/03610927408827101
- [4] Kevin Crowston, Kangning Wei, James Howison, and Andrea Wiggins. 2012. Free/Libre open-source software development: What we know and what we do not know. *ACM Comput. Surv.* 44, 2 (2012), 7:1–7:35. doi:10.1145/2089125.2089127
- [5] Bert J. Dempsey, Debra Weiss, Paul Jones, and Jane Greenberg. 2002. Who is an open source software developer? *Commun. ACM* 45, 2 (2002), 67–72. doi:10.1145/503124.503125
- [6] Leonardo B. Furtado, Bruno Cartaxo, Christoph Treude, and Gustavo Pinto. 2021. How Successful Are Open Source Contributions From Countries With Different Levels of Human Development? *IEEE Softw.* 38, 2 (2021), 58–63. doi:10.1109/MS.2020.3044020
- [7] Jesús M. González-Barahona, Gregorio Robles, Roberto Andradás-Izquierdo, and Rishab Aiyer Ghosh. 2008. Geographic origin of libre software developers. *Inf. Econ. Policy* 20, 4 (2008), 356–363. doi:10.1016/J.INFOECOPOL.2008.07.001
- [8] Aditya Grover and Jure Leskovec. 2016. node2vec: Scalable Feature Learning for Networks. *CoRR abs/1607.00653* (2016). arXiv:1607.00653 <http://arxiv.org/abs/1607.00653>
- [9] Mark S. Handcock, David R. Hunter, Carter T. Butts, Steven M. Goodreau, and Martina Morris. 2008. statnet: Software Tools for the Representation, Visualization, Analysis and Simulation of Network Data. *Journal of Statistical Software* 24, 1 (2008), 1–11. doi:10.18637/jss.v024.i01
- [10] Brandon Heller, Eli Marschner, Evan Rosenfeld, and Jeffrey Heer. 2011. Visualizing collaboration and influence in the open-source software community. In *Proceedings of the 8th International Working Conference on Mining Software Repositories, MSR 2011 (Co-located with ICSE)*, Waikiki, Honolulu, HI, USA, May 21–28, 2011, Proceedings, Arie van Deursen, Tao Xie, and Thomas Zimmermann (Eds.). ACM, 223–226. doi:10.1145/1985441.1985476
- [11] David R. Hunter, Mark S. Handcock, Carter T. Butts, Steven M. Goodreau, and Martina Morris. 2008. ergm: A Package to Fit, Simulate and Diagnose Exponential-Family Models for Networks. *Journal of Statistical Software* 24, 3 (2008), 1–29. doi:10.18637/jss.v024.i03
- [12] Samuel P. Huntington. 1993. The Clash of Civilizations? *Foreign Affairs* 72, 3 (1993), 22–49. <http://www.jstor.org/stable/20045621>
- [13] Matthew O Jackson and Stephen Nei. 2015. Networks of military alliances, wars, and international trade. *Proceedings of the National Academy of Sciences* 112, 50 (2015), 15277–15284.
- [14] Miller McPherson, Lynn Smith-Lovin, and James M Cook. 2001. Birds of a Feather: Homophily in Social Networks. *Annual Review of Sociology* 27, Volume 27, 2001 (2001), 415–444. doi:10.1146/annurev.soc.27.1.415
- [15] Scott Page. 2007. The Difference: How the Power of Diversity Creates Better Groups, Firms, Schools, and Societies. *The Difference: How the Power of Diversity Creates Better Groups, Firms, Schools, and Societies (New Edition)* (01 2007).
- [16] Ayushi Rastogi, Nachiappan Nagappan, and Georgios Gousios. 2016. *Geographical bias in GitHub: Perceptions and reality*. Technical Report. <https://repository.iiitd.edu.in/jspui/bitstream/handle/123456789/388/IIITD-TR-2016-001.pdf?sequence=3>
- [17] Ayushi Rastogi, Nachiappan Nagappan, Georgios Gousios, and André van der Hoek. 2018. Relationship between geographical location and evaluation of developer contributions in GitHub. In *Proceedings of the 12th ACM/IEEE International Symposium on Empirical Software Engineering and Measurement, ESEM 2018, Oulu, Finland, October 11–12, 2018*, Markku Oivo, Daniel Méndez Fernández, and Audris Mockus (Eds.). ACM, 22:1–22:8. doi:10.1145/3239235.3240504
- [18] Garry Robins, Pip Pattison, Yuval Kalish, and Dean Lusher. 2007. An introduction to exponential random graph (p*) models for social networks. *Social networks* 29, 2 (2007), 173–191.
- [19] Davide Rossi and Stefano Zacchiroli. 2022. Geographic Diversity in Public Code Contributions: An Exploratory Large-Scale Study Over 50 Years. In *19th IEEE/ACM International Conference on Mining Software Repositories, MSR 2022, Pittsburgh, PA, USA, May 23–24, 2022*. ACM, 80–85. doi:10.1145/3524842.3528471
- [20] M. Ángeles Serrano, Marián Boguñá, and Alessandro Vespignani. 2009. Extracting the multiscale backbone of complex weighted networks. *Proceedings of the National Academy of Sciences* 106, 16 (April 2009), 6483–6488. doi:10.1073/pnas.0808904106
- [21] David Snyder and Edward L. Kick. 1979. Structural Position in the World System and Economic Growth, 1955–1970: A Multiple-Network Analysis of Transnational Interactions. *Amer. J. Sociology* 84, 5 (1979), 1096–1126. doi:10.1086/226902 arXiv:https://doi.org/10.1086/226902
- [22] Bogdan State, Patrick Park, Ingmar Weber, and Michael Macy. 2015. The Mesh of Civilizations in the Global Network of Digital Communication. *PLOS ONE* 10, 5 (05 2015), 1–9. doi:10.1371/journal.pone.0122543
- [23] Yuri Takhteyev and Andrew Hilts. 2010. Investigating the geography of open source software through GitHub. *Manuscript Submitted for Publication* (2010). <https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=9ba29373aac78aa592f3cbf932fbd2d14d6fcb53>
- [24] Babu Veeresh Thummadi and Srikanth Paruchuri. 2022. Presence of Location-Based Agglomeration Effects in Open Source Communities: An Empirical Test on GitHub. *Academy of Management Discoveries* 8, 2 (2022), 274–297. <https://doi.org/10.5465/amd.2019.0255>
- [25] Sebastian von Engelhardt, Andreas Freytag, and Christoph Schulz. 2013. On the Geographic Allocation of Open Source Software Activities. *Int. J. Innov. Digit. Econ.* 4, 2 (2013), 25–39. doi:10.4018/JIDE.2013040103
- [26] Johannes Wachs, Mariusz Nitecki, William Schueller, and Axel Polleres. 2022. The geography of open source software: Evidence from GitHub. *Technological Forecasting and Social Change* 176 (2022), 121478. <https://doi.org/10.1016/j.techfore.2022.121478>

[27] Stanley Wasserman and Katherine Faust. 1994. *Social Network Analysis: Methods and Applications*. Cambridge University Press.

Received 23 May 2025

1045
1046
1047
1048
1049
1050
1051
1052
1053
1054
1055
1056
1057
1058
1059
1060
1061
1062
1063
1064
1065
1066
1067
1068
1069
1070
1071
1072
1073
1074
1075
1076
1077
1078
1079
1080
1081
1082
1083
1084
1085
1086
1087
1088
1089
1090
1091
1092
1093
1094
1095
1096
1097
1098
1099
1100
1101
1102

1103
1104
1105
1106
1107
1108
1109
1110
1111
1112
1113
1114
1115
1116
1117
1118
1119
1120
1121
1122
1123
1124
1125
1126
1127
1128
1129
1130
1131
1132
1133
1134
1135
1136
1137
1138
1139
1140
1141
1142
1143
1144
1145
1146
1147
1148
1149
1150
1151
1152
1153
1154
1155
1156
1157
1158
1159
1160